

# Data Race Example

```
public class Example extends Thread {  
    private static int cnt = 0;    // shared state  
    public void run() {  
        int y = cnt;  
        cnt = y + 1;  
    }  
    public static void main(String args[]) {  
        Thread t1 = new Example();  
        Thread t2 = new Example();  
        t1.start();  
        t2.start();  
    }  
}
```

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 0

```
t1.run() {  
    int y = cnt;  
    cnt = y + 1;  
}
```

```
t2.run() {  
    int y = cnt;  
    cnt = y + 1;  
}
```

*Start: both threads ready to run. Each will increment the global count.*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 0

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;
```

```
    cnt = y + 1;
```

```
}
```

*T1 executes, grabbing  
the global counter value into y.*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 0

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```

*T1 is pre-empted. T2  
executes, grabbing the global  
counter value into y.*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    **cnt = 1**

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```

*T2 executes, storing the incremented cnt value.*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    **cnt = 1**

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```

*T2 completes. T1  
executes again, storing the  
old counter value (1) rather  
than the new one (2)!*

But When I Run it Again?

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 0

```
t1.run() {  
    int y = cnt;  
    cnt = y + 1;  
}
```

```
t2.run() {  
    int y = cnt;  
    cnt = y + 1;  
}
```

*Start: both threads ready to run. Each will increment the global count.*



# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 0

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;
```

```
    cnt = y + 1;
```

```
}
```

*T1 executes, grabbing  
the global counter value into y.*

# Data Race Example

```
static int cnt = 0;
t1.run() {
    int y = cnt;
    cnt = y + 1;
}
t2.run() {
    int y = cnt;
    cnt = y + 1;
}
```

*Shared state*    **cnt = 1**

*y = 0*



*T1 executes again, storing the counter value*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    cnt = 1

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;    y = 1
```

```
    cnt = y + 1;
```

```
}
```

*T1 finishes. T2 executes,  
grabbing the global  
counter value into y.*

# Data Race Example

```
static int cnt = 0;
```

*Shared state*    **cnt = 2**

```
t1.run() {
```

```
    int y = cnt;    y = 0
```

```
    cnt = y + 1;
```

```
}
```



```
t2.run() {
```

```
    int y = cnt;    y = 1
```

```
    cnt = y + 1;
```

```
}
```

*T2 executes, storing the incremented cnt value.*

# What Happened?

- In the first example, **t1** was preempted after it read the counter but before it stored the new value.
  - Depends on the idea of an *atomic action*
  - Violated an object invariant
- A particular way in which the execution of two threads is interleaved is called a *schedule*. We want to prevent this undesirable schedule.
- Undesirable schedules can be hard to reproduce, and so hard to debug.

# Question

- If instead of

```
int y = cnt;  
cnt = y+1;
```
- We had written

```
– cnt++;
```
- Would the result be any different?
- Answer: NO!
  - Don't depend on your intuition about atomicity

# Question

- If you run a program with a race condition, will you always get an unexpected result?
  - No! It depends on the scheduler
  - ...i.e., which JVM you're running
  - ...and on the other threads/processes/etc that are running on the same CPU
- Race conditions are hard to find

# What can this example generate?

```
static int cnt = 2;
t1.run() {
    int y = cnt;
    cnt = y * 2;
}
t2.run() {
    int y = cnt;
    cnt = y + 1;
}
```



# Avoiding Interference: Synchronization

```
public class Example extends Thread {  
    private static int cnt = 0;  
    static Object lock = new Object();  
    public void run() {  
        synchronized (lock) {  
            int y = cnt;  
            cnt = y + 1;  
        }  
    }  
    ...  
}
```

*Lock, for protecting  
the shared state*

*Acquires the lock;  
Only succeeds if not  
held by another  
thread*

*Releases the lock*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    cnt = 0



*T1 acquires the lock*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    cnt = 0



*T1 reads cnt into y*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    cnt = 0



*T1 is pre-empted.  
T2 attempts to  
acquire the lock but fails  
because it's held by  
T1, so it blocks*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    **cnt = 1**



*T1 runs, assigning  
to cnt*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    cnt = 1



*T1 releases the lock  
and terminates*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;
    }
}
```

*Shared state*    cnt = 1



*T2 now can acquire  
the lock.*

# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 1
    }
}
```

*Shared state*    cnt = 1



*T2 reads cnt into y.*



# Applying Synchronization

```
int cnt = 0;
t1.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 0
    }
}
t2.run() {
    synchronized(lock) {
        int y = cnt;
        cnt = y + 1;    y = 1
    }
}
```

*Shared state*    **cnt = 2**



*T2 assigns cnt,  
then releases the lock*

# Locks

- *Any* Object subclass has (can act as) a lock
- Only one thread can hold the lock on an object
  - Other threads block until they can acquire it
- If a thread already holds the lock on an object
  - The thread can reacquire the same lock many times
    - ...Locks are *reentrant*
  - Lock is released when object unlocked the corresponding number of times
- No way to attempt to acquire a lock in Java 1.4
  - Either succeeds, or blocks the thread

# Synchronized Statement

- **synchronized (obj) { *statements* }**
- Obtains the lock on **obj** before executing statements in block
- Releases the lock when the statement block completes
  - Either normally, or due to a return, break, or exception being thrown in the block

# Synchronized Methods

- A method can be synchronized
  - Add **synchronized** modifier before return type
- Obtains the lock on object referenced by **this** before executing method
  - Releases lock when method completes
- For a **static synchronized** method
  - Locks the **Class** object for the class
    - Accessible directly, e.g. **Foo.class**
  - Not the same as **this** (there is no this)

# Synchronization Example

```
public class State {  
    private int cnt = 0;  
    public synchronized void incCnt(int x) {  
        cnt += x;  
    }  
    public synchronized int getCnt() { return cnt; }  
}  
public class MyThread extends Thread {  
    State s;  
    public MyThread(State s) { this.s = s; }  
    public void run() {  
        s.incCnt(1)  
    }  
    public void main(String args[]) {  
        State s = new State();  
        MyThread thread1 = new MyThread(s);  
        MyThread thread2 = new MyThread(s);  
        thread1.start(); thread2.start();  
    }  
}
```

*Synchronization  
occurs in State  
object itself,  
rather than in  
its caller.*

# Locking

- The object locked, and the fields protected by the lock, may be different
- It is common to protect fields by holding a lock on the object they are fields of
  - but this is not required
- For example, fields of all the nodes of a tree might be protected by a lock on the tree

# Synchronization Style

- Design decision
  - Internal synchronization (class is thread-safe)
    - Have a stateful object synchronize itself (e.g., with synchronized methods)
  - External synchronization (class is thread-compatible)
    - Have callers perform synchronization before calling the object
- Can go both ways:
  - Thread-safe: Random
  - Thread-compatible: ArrayList, HashMap, ...

# Need for Synchronization

- Need to use synchronization to ensure that two threads cannot access a shared memory location at the same time
  - Unless both are reading the memory
  - **volatile** is a special case we'll come to later



# What can go wrong?

- Deadlock
- Insufficient atomicity
- Non-determinism
- Just plain wrong

# Synchronization not a Panacea

- Two threads can block on locks held by the other; this is called *deadlock*

```
Object A = new Object();
```

```
Object B = new Object();
```


```
T1.run() {  
    synchronized (A) {  
        synchronized (B) {  
            ...  
        }  
    }  
}
```

```
T2.run() {  
    synchronized (B) {  
        synchronized (A) {  
            ...  
        }  
    }  
}
```


# Deadlock

- Quite possible to create code that deadlocks
  - Thread 1 holds lock on **A**
  - Thread 2 holds lock on **B**
  - Thread 1 is trying to acquire a lock on **B**
  - Thread 2 is trying to acquire a lock on **A**
  - Deadlock!
- Not easy to detect when deadlock has occurred
  - Other than by the fact that nothing is happening

# Deadlock: Wait graphs




Thread T1 holds lock A



Thread T2 attempting to acquire lock B

Deadlock occurs when there is a cycle in the graph

# Wait graph example



T1 holds lock on **A**

T2 holds lock on **B**

T1 is trying to acquire a lock on **B**

T2 is trying to acquire a lock on **A**

# Deadlock happens

- Consider:
  - each bank account is protected by a distinct lock
  - Operations on an account acquire the lock on that account
  - You want to move money from one account to the other
    - You need to hold both locks

# Avoiding deadlock

- Can try to never hold a lock on more than one object at a time
  - Be careful what you call while holding a lock
- Can impose a lock ordering
  - A total order such that if  $X < Y$  and I hold a lock on both  $X$  and  $Y$ , I grab the lock on  $X$  before the lock on  $Y$

# Insufficient Atomicity

```
public class State {
    private int cnt = 0;
    public synchronized int getCnt() {
        return cnt;
    }
    public synchronized void setCnt(int newValue) {
        cnt = newValue;
    }
}

public class MyThread extends Thread {
    State s;
    public MyThread(State s) { this.s = s; }
    public void run() {
        s.setCnt(s.getCnt()+1);
    }
}

public void main(String args[]) {
    State s = new State();
    MyThread thread1 = new MyThread(s);
    MyThread thread2 = new MyThread(s);
    thread1.start(); thread2.start();
}
}
```



# Insufficient Atomicity

- In a number of situations, you will want to
  - Examine the current state of the system
  - Update the system based on that examination
- You write this assuming that this will be done as an atomic action
  - But if the operations to examine the system and update the system are separate synchronized methods, doesn't work

# Synchronized Maps

```
public class WordCounter extends Thread {  
  
    static Map<String, Integer> wordCount  
        = Collections.synchronizedMap(  
            new HashMap<String, Integer>());  
  
    public void run() {  
        while (...) {  
            String word = ...;  
            Integer v = wordCount.get(word);  
            if (v == null) wordCount.put(word, 1);  
            else wordCount.put(word, v+1);  
        }  
    }  
}
```

# Nondeterminism

- Even if your program is entirely “correct”, it may be non-deterministic
  - Not consistently producing the same result
- Makes testing your code difficult
- Makes finding bugs difficult
- Generally, no fix for nondeterminism in multithreaded code

# Just plain wrong

- Multithreaded code can be subtle, and testing is difficult
- Easy to make mistakes when writing multithreaded code
- Use building blocks built by experts if possible
  - [java.util.concurrent](#)